

# CS315: Principles of Database Systems

## Database Transactions

Arnab Bhattacharya  
`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur  
<http://web.cse.iitk.ac.in/~cs315/>

2<sup>nd</sup> semester, 2013-14  
Tue, Fri 1530-1700 at CS101

# ACID properties

- A **transaction** is a logical unit of a program
- To preserve data integrity, a database must follow four properties
  - 1 **Atomicity**: Either all operations of a transaction are reflected or none are reflected
  - 2 **Consistency**: If a database is consistent before the execution of the transaction, it must be consistent after it
  - 3 **Isolation**: Although multiple transactions may execute concurrently, each transaction must be unaware of others, i.e., to a transaction, it must seem that either any other transaction has completed execution or has not started execution at all
  - 4 **Durability**: After a transaction finishes successfully, the changes must be permanent in the database despite subsequent failures
- Together, these four properties are called the **ACID** properties

# Transaction model

- A **transaction** is a logical unit of a program that reads and writes various named data items
- Thus, a database is simply a collection of *named items*
- *Granularity* of data: field, record, block, relation
- Transaction concepts are independent of granularity

# Transaction model

- A **transaction** is a logical unit of a program that reads and writes various named data items
- Thus, a database is simply a collection of *named items*
- *Granularity* of data: field, record, block, relation
- Transaction concepts are independent of granularity
- Two basic operations: **read** and **write**
- **read(x)**: reads a database item named x
- **write(x)**: writes to a database item named x
- Database program uses the same name x for the variable

# Transaction model

- A **transaction** is a logical unit of a program that reads and writes various named data items
- Thus, a database is simply a collection of *named items*
- *Granularity* of data: field, record, block, relation
- Transaction concepts are independent of granularity
- Two basic operations: **read** and **write**
- **read(x)**: reads a database item named x
- **write(x)**: writes to a database item named x
- Database program uses the same name x for the variable
- A transaction can read or write a data item only once
- Conflicts with read and write

# Transaction model

- A **transaction** is a logical unit of a program that reads and writes various named data items
- Thus, a database is simply a collection of *named items*
- *Granularity* of data: field, record, block, relation
- Transaction concepts are independent of granularity
- Two basic operations: **read** and **write**
- **read(x)**: reads a database item named x
- **write(x)**: writes to a database item named x
- Database program uses the same name x for the variable
- A transaction can read or write a data item only once
- Conflicts with read and write
  - **RAW**: read-after-write
  - **WAR**: write-after-read
  - **WAW**: write-after-write

# Example

- Transfer 50 rupees from account  $a$  to account  $b$
- `read(a); a := a - 50; write(a); read(b); b := b + 50; write(b)`
- Atomicity:

# Example

- Transfer 50 rupees from account  $a$  to account  $b$
- `read(a); a := a - 50; write(a); read(b); b := b + 50; write(b)`
- Atomicity:  $a$  must not be debited without crediting  $b$
- Consistency:



# Example

- Transfer 50 rupees from account  $a$  to account  $b$
- `read(a); a := a - 50; write(a); read(b); b := b + 50; write(b)`
- Atomicity:  $a$  must not be debited without crediting  $b$
- Consistency: Sum of  $a$  and  $b$  remains constant
- Isolation:

# Example

- Transfer 50 rupees from account  $a$  to account  $b$
- `read(a); a := a - 50; write(a); read(b); b := b + 50; write(b)`
- Atomicity:  $a$  must not be debited without crediting  $b$
- Consistency: Sum of  $a$  and  $b$  remains constant
- Isolation: If another transaction reads  $a$  and  $b$ , it must see one of the two states – before or after the transaction, otherwise  $a + b$  is wrong
- Durability:

# Example

- Transfer 50 rupees from account  $a$  to account  $b$
- `read(a); a := a - 50; write(a); read(b); b := b + 50; write(b)`
- Atomicity:  $a$  must not be debited without crediting  $b$
- Consistency: Sum of  $a$  and  $b$  remains constant
- Isolation: If another transaction reads  $a$  and  $b$ , it must see one of the two states – before or after the transaction, otherwise  $a + b$  is wrong
- Durability: If  $b$  is notified of credit, it must persist even if the database crashes

# Problems in transactions

- **Lost update:** Update by one transaction is overwritten by another transaction

# Problems in transactions

- **Lost update**: Update by one transaction is overwritten by another transaction
- **Temporary update** or **Dirty read**: A transaction updates and then fails; another transaction reads it before it is reverted to original value

# Problems in transactions

- **Lost update**: Update by one transaction is overwritten by another transaction
- **Temporary update** or **Dirty read**: A transaction updates and then fails; another transaction reads it before it is reverted to original value
- **Incorrect summary**: One transaction is updating values while other is computing an aggregate on them

# Causes of transaction failures

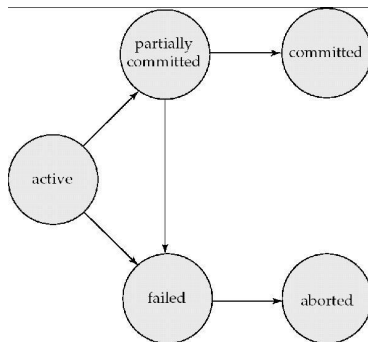
# Causes of transaction failures

- System crash
  - Memory is lost
- System error
  - Divide by zero
- Exceptions
  - Insufficient account balance
- Concurrency enforcement
  - Deadlock detection
- Disk crash
  - Persistency fails
- Physical problems
  - Power failure, fire



# Transaction state

- **Active**: transaction is executing
- **Partially committed**: after last statement has been executed
- **Failed**: when execution cannot proceed
- **Committed**: after successful completion
- **Aborted**: transaction has been rolled back and it has been ensured that there is no effect of the transaction



# Shadow database scheme

- **Recovery management system** of a database ensures that atomicity and durability properties are maintained
- **Shadow database** scheme enforces atomicity
  - A *shadow copy* of the database is made before any transaction
  - All updates are made on the shadow copy
  - If the transaction finishes successfully, the database pointer is updated to the shadow copy
  - If the transaction fails, the old database pointer which points to the original database is maintained
- Inefficient for large databases
- Cannot efficiently handle concurrent transactions

- **Log** or **journal** keeps track of all transaction operations
- Log enforces atomicity and durability
  - Log is maintained on disk
  - Log is periodically backed up to archival storage to guard against disk failures
- For a transaction  $T$ , following log records are maintained
  - (start,  $T$ )
  - (write,  $T$ ,  $x$ , old, new)
  - (read,  $T$ ,  $x$ , val)
  - (commit,  $T$ )
  - (abort,  $T$ )

- **Log** or **journal** keeps track of all transaction operations
- Log enforces atomicity and durability
  - Log is maintained on disk
  - Log is periodically backed up to archival storage to guard against disk failures
- For a transaction  $T$ , following log records are maintained
  - (start,  $T$ )
  - (write,  $T$ ,  $x$ , old, new)
  - (read,  $T$ ,  $x$ , val)
  - (commit,  $T$ )
  - (abort,  $T$ )
- Recovery using logs
  - **Undo** possible by traversing log *backward* and setting database items to *old* values
  - **Redo** possible by traversing log *forward* and setting database items to *new* values

# Commit point

- A transaction reaches its **commit point** when all its operations have been executed successfully and they have been recorded in the log
- Beyond the commit point, the transaction is said to be *committed*, and its effect on the database is assumed to be permanent
- A *commit entry* (commit, T) is made in the log

# Commit point

- A transaction reaches its **commit point** when all its operations have been executed successfully and they have been recorded in the log
- Beyond the commit point, the transaction is said to be *committed*, and its effect on the database is assumed to be permanent
- A *commit entry* (commit, T) is made in the log
- A transaction needs to **rollback** if there is a (start, T) entry but no (commit, T) entry
  - It may also need to be rolled back if there is a (abort, T) entry
  - Undo operations need to be performed

# Commit point

- A transaction reaches its **commit point** when all its operations have been executed successfully and they have been recorded in the log
- Beyond the commit point, the transaction is said to be *committed*, and its effect on the database is assumed to be permanent
- A *commit entry* (commit, T) is made in the log
- A transaction needs to **rollback** if there is a (start, T) entry but no (commit, T) entry
  - It may also need to be rolled back if there is a (abort, T) entry
  - Undo operations need to be performed
- Before a transaction reaches its commit point, any portion of the log not yet written to disk must be flushed – this is called **force-writing** of the log
  - Ensures that redo operations can be done successfully

# Concurrency

- Multiple transactions should be able to run concurrently
- Advantages
  - Increased processor and disk utilization leading to better *throughput*: one is using CPU, other is doing disk I/O
  - Reduced average response time: short transactions finish earlier and do not wait behind long ones
- **Concurrency control schemes** achieve isolation
- Must ensure **correctness** of concurrent executions
- **Serializability** imposes notion of correctness