

CS315: Principles of Database Systems

Schedules

Arnab Bhattacharya
`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,
Indian Institute of Technology, Kanpur
<http://web.cse.iitk.ac.in/~cs315/>

2nd semester, 2013-14
Tue, Fri 1530-1700 at CS101

Schedule

- A **schedule** is a chronological sequence of instructions from concurrent transactions
- If a transaction appears in a schedule, *all* instructions of the transaction must appear in the schedule
- Order of instructions within a transaction must be maintained in the schedule
- A transaction finishing successfully will have *commit* as the last instruction
- A transaction not finishing successfully will have *abort* as the last instruction
- Commit and abort statements may be omitted if obvious

Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:

$r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$

$r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$

Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B

- A **serial** schedule:

$r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$

- Another schedule:

$r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$

Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$
- Another schedule:
 $r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$
- This is not a serial schedule but is **equivalent**

Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B

- A **serial** schedule:

$r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$

- Another schedule:

$r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$

- This is not a serial schedule but is **equivalent**

- Yet another schedule:

$r_1(A); A := A - 50; r_2(A); t := 0.1A; A := A - t; w_2(A);$
 $w_1(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$

Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$
- Another schedule:
 $r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$
- This is not a serial schedule but is **equivalent**
- Yet another schedule:
 $r_1(A); A := A - 50; r_2(A); t := 0.1A; A := A - t; w_2(A);$
 $w_1(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$
- This is not a serial schedule and is not equivalent either

Serializability

- Each transaction preserves database consistency
- Hence, a serial schedule also does that
- A schedule is **serializable** if it is *equivalent* to a serial schedule
- There are different forms of equivalence giving rise to notions of
 - **Conflict serializability**
 - **View serializability**
- Operations other than *read* and *write* are ignored

Serializability

- Each transaction preserves database consistency
- Hence, a serial schedule also does that
- A schedule is **serializable** if it is *equivalent* to a serial schedule
- There are different forms of equivalence giving rise to notions of
 - **Conflict serializability**
 - **View serializability**
- Operations other than *read* and *write* are ignored
- Instruction I_i of transaction T_i **conflicts** with I_j of T_j if and only if they access the **same** data item and at least one of them is a **write**
- Intuitively, a conflict enforces a logical temporal order of the instructions
- Consequently, if two instructions do not conflict, they can be interchanged

Conflict serializability

- A schedule S is **conflict equivalent** to another schedule S' if it can be transformed to S' by a series of swaps of *non-conflicting* instructions
- A schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule
- A serial schedule is conflict serializable, but not vice versa
- If a schedule is conflict serializable, it is correct in the sense that it preserves database consistency

Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$

Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$
is conflict serializable as it is conflict equivalent to the serial schedule
 $T_1T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
 - It is not required to be conflict equivalent to T_2T_1 as well
- $r_1(a)w_2(a)w_1(a)$

Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$
is conflict serializable as it is conflict equivalent to the serial schedule $T_1T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
 - It is not required to be conflict equivalent to T_2T_1 as well
- $r_1(a)w_2(a)w_1(a)$
is *not* conflict serializable as it is not conflict equivalent to either of the two serial schedules T_1T_2 and T_2T_1

View serializability

- Two schedules are **view equivalent** if the reads in them get the same “**view**”, i.e., they read the value produced by the same write operation

View serializability

- Two schedules are **view equivalent** if the reads in them get the same “**view**”, i.e., they read the value produced by the same write operation
- Formally, two schedules S and S' are **view equivalent** if
 - 1 For each data item x , if a transaction T reads the initial value of x in S , it reads the same initial value of x in S' as well
 - 2 For each data item x , if a transaction T writes the final value of x in S , it writes the final value of x in S' as well
 - 3 If transaction T_i reads the value of data item x produced by write by transaction T_j in S , it must read the value written by T_j in S' as well
- A schedule S is **view serializable** if it is view equivalent to a serial schedule

Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$

Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$
is view serializable as it is view equivalent to the serial schedule
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
- $r_1(a)w_2(a)w_1(a)w_3(a)$

Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$
is view serializable as it is view equivalent to the serial schedule
 $T_1T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
- $r_1(a)w_2(a)w_1(a)w_3(a)$
is view serializable as it is view equivalent to the serial schedule
 $T_1T_2T_3 : r_1(a)w_1(a)w_2(a)w_3(a)$
- $r_1(a)w_2(a)w_1(a)$

Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$
is view serializable as it is view equivalent to the serial schedule
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
- $r_1(a)w_2(a)w_1(a)w_3(a)$
is view serializable as it is view equivalent to the serial schedule
 $T_1 T_2 T_3 : r_1(a)w_1(a)w_2(a)w_3(a)$
- $r_1(a)w_2(a)w_1(a)$
is *not* view serializable as it is not view equivalent to either of the two
serial schedules $T_1 T_2$ and $T_2 T_1$

Relationship between view and conflict serializability

Relationship between view and conflict serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability

Relationship between view and conflict serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability
- They are same under the **constrained write assumption**
- In this assumption, every write of a data item x is constrained by the value of x it has read
 - $\text{write}(f(\text{read}(x)))$

Relationship between view and conflict serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability
- They are same under the **constrained write assumption**
- In this assumption, every write of a data item x is constrained by the value of x it has read
 - $\text{write}(f(\text{read}(x)))$
- With unconstrained writes (**blind writes**), a schedule that is view serializable is not necessarily conflict serializable
 - Blind writes: writes to a data item without reading it

Relationship between view and conflict serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability
- They are same under the **constrained write assumption**
- In this assumption, every write of a data item x is constrained by the value of x it has read
 - $\text{write}(f(\text{read}(x)))$
- With unconstrained writes (**blind writes**), a schedule that is view serializable is not necessarily conflict serializable
 - Blind writes: writes to a data item without reading it
- Every view serializable schedule that is not conflict serializable must have blind writes

Other notions of equivalence

- Conflict and view serializable schedules are restrictive in the sense that they aim to guarantee database consistency without analyzing the result

Other notions of equivalence

- Conflict and view serializable schedules are restrictive in the sense that they aim to guarantee database consistency without analyzing the result
- A schedule S is **result equivalent** to a schedule S' if it produces the same result as S'
- Consider
$$r_1(A); A := A - 50; w_1(A); r_2(B); B := B - 10; w_2(B);$$
$$r_1(B); B := B + 50; w_1(B); r_2(A); A := A + 10; w_2(A);$$
- It produces the **same result** as the serial schedule
$$r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$$
$$r_2(B); B := B - 10; w_2(B); r_2(A); A := A + 10; w_2(A);$$
but is not conflict or view serializable
- Determining such equivalence requires *semantic* analysis of operations other than read and write

Testing for serializability

- Create a **precedence graph** for the schedule
- Directed graph where each transaction is a vertex
- An edge from transaction T_i to T_j exists if
 - $w_i(x)$ precedes $r_j(x)$, or
 - $r_i(x)$ precedes $w_j(x)$, or
 - $w_i(x)$ precedes $w_j(x)$
- A schedule is conflict serializable if and only if its precedence graph is *acyclic*
- *Depth-first search* can detect cycles in $O(n + m)$ time
- *Topological sorting* produces an equivalent serial order

Testing for serializability

- Create a **precedence graph** for the schedule
- Directed graph where each transaction is a vertex
- An edge from transaction T_i to T_j exists if
 - $w_i(x)$ precedes $r_j(x)$, or
 - $r_i(x)$ precedes $w_j(x)$, or
 - $w_i(x)$ precedes $w_j(x)$
- A schedule is conflict serializable if and only if its precedence graph is *acyclic*
- *Depth-first search* can detect cycles in $O(n + m)$ time
- *Topological sorting* produces an equivalent serial order
- Testing for view serializability is *NP-complete*
- Practical algorithms
 - Catches all non view serializable schedules
 - But can miss a view serializable schedule

Recoverable schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important

Recoverable schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important
- A schedule is called a **recoverable schedule** if
 - A transaction T_i reads a data item previously written by T_j , and
 - T_j commits *before* T_i commits
- Consider $r_1(a)w_1(a)r_2(a)r_1(b)$

Recoverable schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important
- A schedule is called a **recoverable schedule** if
 - A transaction T_i reads a data item previously written by T_j , and
 - T_j commits *before* T_i commits
- Consider $r_1(a)w_1(a)r_2(a)r_1(b)$
- If T_2 commits just after $r_2(a)$, i.e., if the schedule is $r_1(a)w_1(a)r_2(a)c_2r_1(b)a_1$, then it is *not* recoverable
 - If T_1 crashes, then $w_1(a)$ is undone, but T_2 has already read a wrong value of a and committed

Recoverable schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important
- A schedule is called a **recoverable schedule** if
 - A transaction T_i reads a data item previously written by T_j , and
 - T_j commits *before* T_i commits
- Consider $r_1(a)w_1(a)r_2(a)r_1(b)$
- If T_2 commits just after $r_2(a)$, i.e., if the schedule is $r_1(a)w_1(a)r_2(a)c_2r_1(b)a_1$, then it is *not* recoverable
 - If T_1 crashes, then $w_1(a)$ is undone, but T_2 has already read a wrong value of a and committed
- Therefore, to make it recoverable, the schedule should be $r_1(a)w_1(a)r_2(a)r_1(b)c_1c_2$
 - If T_1 aborts, T_2 can also abort

Cascading rollbacks

- In recoverable schedules, a single transaction failure may lead to a series of rollbacks
- This is called **cascading rollbacks** or **cascading aborts**
- Consider $r_1(a)w_1(a)r_2(a)w_2(a)r_3(a)r_1(b)a_1c_2c_3$
- It is recoverable
- However, if T_1 fails, T_2 and T_3 must abort as well
- Not preferable as lot of work is undone

Cascadeless schedule

- Cascading rollbacks are eliminated

Cascadeless schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
 - A transaction T_i reads a data item previously written by T_j , and
 - T_j commits before T_i reads
- Consider $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$

Cascadeless schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
 - A transaction T_i reads a data item previously written by T_j , and
 - T_j commits before T_i reads
- Consider $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$
- It is not cascadeless as T_2 reads a written by T_1 before T_1 commits

Cascadeless schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
 - A transaction T_i reads a data item previously written by T_j , and
 - T_j commits before T_i reads
- Consider $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$
- It is not cascadeless as T_2 reads a written by T_1 before T_1 commits
- Therefore, to make it cascadeless, the schedule should be $r_1(a)w_1(a)r_1(b)c_1r_2(a)c_2$
- No completed transaction needs to be rolled back

Cascadeless schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
 - A transaction T_i reads a data item previously written by T_j , and
 - T_j commits before T_i reads
- Consider $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$
- It is not cascadeless as T_2 reads a written by T_1 before T_1 commits
- Therefore, to make it cascadeless, the schedule should be $r_1(a)w_1(a)r_1(b)c_1r_2(a)c_2$
- No completed transaction needs to be rolled back
- Every cascadeless schedule is recoverable, but not vice versa

Strict schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write

Strict schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
 - A transaction T_i reads or writes a data item previously written by T_j , and
 - T_j commits before T_i reads or writes
- Consider $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$

Strict schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
 - A transaction T_i reads or writes a data item previously written by T_j , and
 - T_j commits before T_i reads or writes
- Consider $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$
- It is not strict as T_2 writes a written by T_1 before T_1 commits

Strict schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
 - A transaction T_i reads or writes a data item previously written by T_j , and
 - T_j commits before T_i reads or writes
- Consider $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$
- It is not strict as T_2 writes a written by T_1 before T_1 commits
- Therefore, to make it strict, the schedule should be $r_1(a)w_1(a)r_1(b)c_1w_2(a)c_2$

Strict schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
 - A transaction T_i reads or writes a data item previously written by T_j , and
 - T_j commits before T_i reads or writes
- Consider $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$
- It is not strict as T_2 writes a written by T_1 before T_1 commits
- Therefore, to make it strict, the schedule should be $r_1(a)w_1(a)r_1(b)c_1w_2(a)c_2$
- Every strict schedule is cascadeless, but not vice versa

Relationship between schedules

